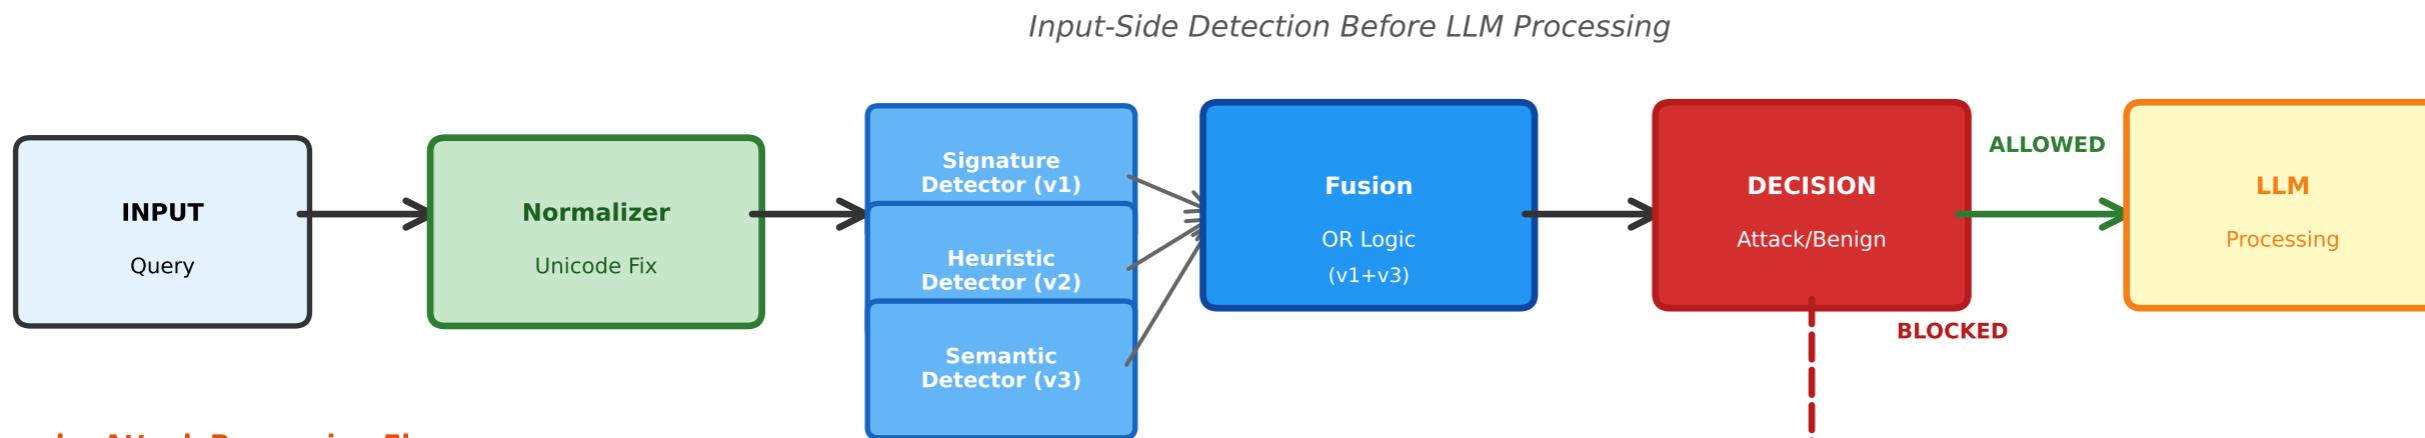
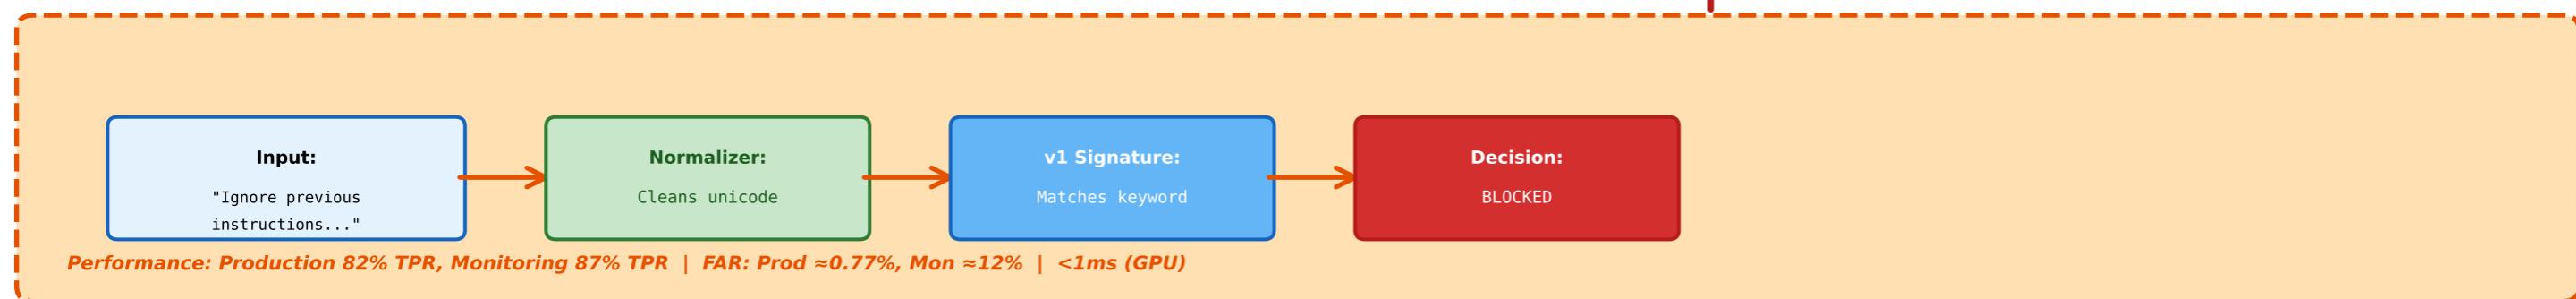


Prompt Injection Detection Pipeline Architecture



Example: Attack Processing Flow



Production Configuration: Normalizer + v3

True Positive Rate (TPR): 82%
False Alarm Rate (FAR): 0.77%
Latency: <1ms per sample (GPU)
Complexity: ~1,200 lines
Deployment: Stateless
Dependencies: sentence-transformers, torch

Component Specifications

Signature Detector (v1):

- 89% TPR, 0% FAR (P1)
- Keyword matching

 Semantic Detector (v3):

- 82% TPR, 0% FAR (P1)
- Pattern analysis

 Fusion: OR Logic (v1+v3)

- Monitoring: 87% TPR

Key Design Principles

1. INPUT-SIDE DETECTION: Attacks blocked BEFORE reaching the LLM
2. NORMALIZER FIRST: Unicode/homoglyph normalization ensures consistent detection
3. COMPLEMENTARY DETECTORS: v1 (signature) + v3 (semantic) catch different patterns
4. THRESHOLD-INVARIANT: Binary OR logic eliminates threshold tuning complexity
5. PRODUCTION-READY: <1ms latency with GPU acceleration, stateless architecture

Legend:

□ Input

□ Normalizer

□ Detector

□ Fusion

□ Decision

□ LLM